

HuanHusted_TSA_Competition

Kelsey Husted & Yu Huan

2023-04-28

Time Series Competition

```
knitr::opts_chunk$set(echo = TRUE, tidy.opts=list(width.cutoff=80), tidy=FALSE)
```

```
##Load packages
```

```
library(readxl)
library(dplyr)
library(lubridate)
library(openxlsx)
library(ggplot2)
library(forecast)
library(Kendall)
library(tseries)
library(outliers)
library(tidyverse)
library(smooth)
library(zoo)
library(kableExtra)
```

Data wrangling and processing from 2005-2009

Hourly data was transformed into daily data with aggregate functions and pipes (i.e., tidyverse).

```
#Import data
df <- read_xlsx("./Competition/Data/load.xlsx")

df_T <- read_xlsx("./Competition/Data/temperature.xlsx")

df_H <- read_xlsx("./Competition/Data/relative_humidity.xlsx")

#Wrangle data from hourly to daily
#Wrangling date column 2005 to 2009
df_daily <- df %>%
  mutate( Date = ymd(date)) %>%
  filter(Date < '2010-01-01')

#removing no numeric columns so rowMeans() function will work
df_dailyV2 <- df %>%
  mutate( Date = ymd(date)) %>%
  filter(Date < '2010-01-01') %>%
  select(3:26)
```

```

#Creating daily data
df_processed <- df_dailyV2 %>%
  mutate(rowMeans(df_dailyV2)) %>%
  rename(Daily_data = "rowMeans(df_dailyV2)") %>%
  select(25)

#Combining date and daily data
date <- df_daily[,2]
df_processed <- cbind(date, df_processed)

nobs = nrow(df_daily)

```

Data wrangling and processing from 2005-2010

The data needs to be formatted to include 2010 as well since the objective is to forecast for 2011. Instead of making two separate datasets, I should use the `window()` function for future reference.

```

#Wrangle data from hourly to daily
#Wrangling date column 2005 to 2010
#removing no numeric columns so rowMeans() function will work
df_daily2010 <- df %>%
  mutate( Date = ymd(date)) %>%
  select(3:26)

#Creating daily data
df_processed2010 <- df_daily2010 %>%
  mutate(rowMeans(df_daily2010)) %>%
  rename(Daily_data = "rowMeans(df_daily2010)") %>%
  select(25)

#Combining data and daily data
date <- df[,2]
df_processed2010 <- cbind(date, df_processed2010)

nobs2010 = nrow(df_processed2010)

```

Time series object transformation

```

#ts transformation 2005 to 2009
ts_daily <- msts(df_processed$Daily_data,
  seasonal.periods=c(7,365.25),
  start=c(2005, 01, 01))


#ts transformation 2005 to 2010
ts_daily2010 <- msts(df_processed2010$Daily_data,
  seasonal.periods=c(7,365.25),
  start=c(2005, 01, 01))

#ts for accuracy test visalization
ts_daily2010_test <- msts(df_processed2010$Daily_data[1:1885],
  seasonal.periods=c(7,365.25),
  start=c(2005, 01, 01))

```

Time series decomposition and plot


```
#Decompose time series  
ts_decompose <- ts_daily %>%  
  mstl()  
plot(ts_decompose)
```




Model 1 STL + ETS: Forecast 2010

```
#Fit and forecast STL model  
ETS_model <- stlf(ts_daily,h=59)  
  
#Plot forecasting  
autoplot(ETS_model) + ylab("Forecast")
```

Forecasts from STL + ETS(A,N,N)



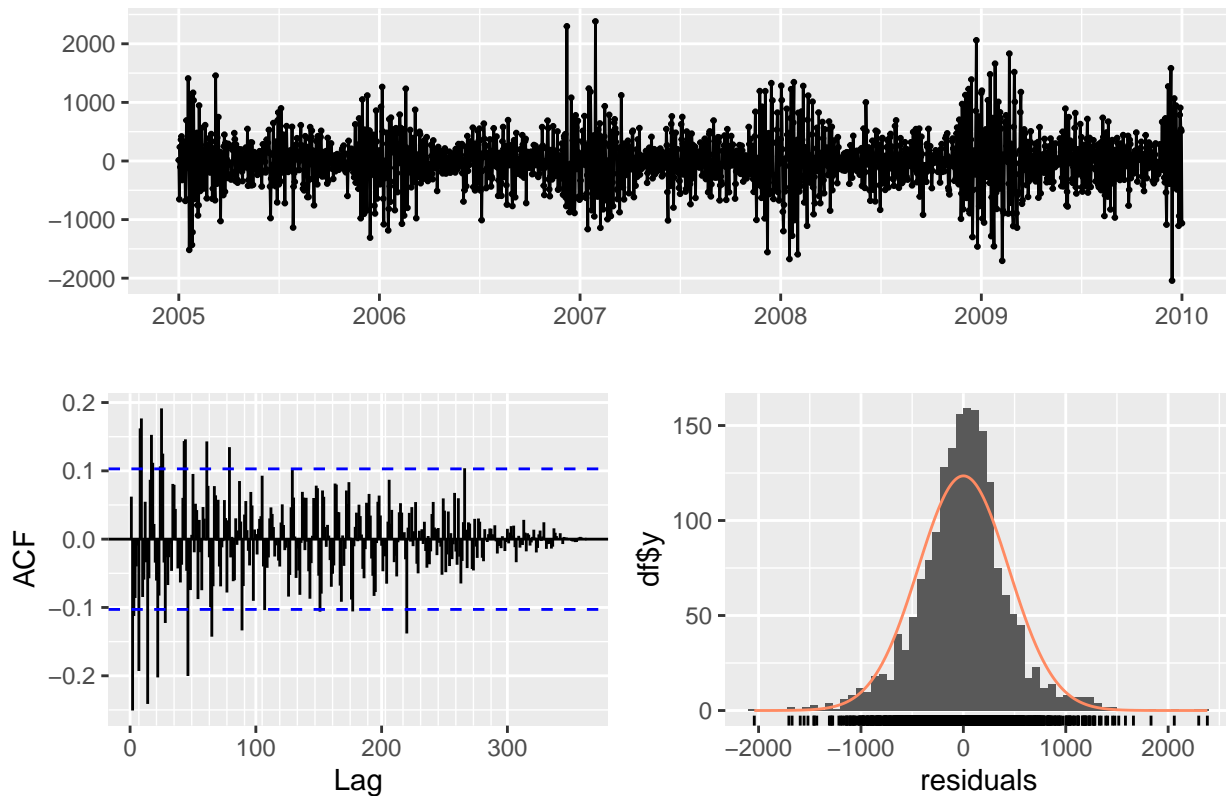
```
#Plot model + observed data  
autoplot(ts_daily2010_test) +  
  autolayer(ETS_model, series="STL Model",PI=FALSE) +  
  ylab("Daily Demand")
```



```
#Plot the residuals
```

```
checkresiduals(ETS_model)
```

Residuals from STL + ETS(A,N,N)



```
##
##  Ljung-Box test
##
## data:  Residuals from STL +  ETS(A,N,N)
## Q* = 1309.7, df = 363, p-value < 2.2e-16
##
## Model df: 2.   Total lags used: 365
#Check accuracy of model
n_for <- 59
observed <- df_processed2010[1827:1885, "Daily_data"]
ETS_scores <- accuracy(ETS_model$mean,observed)
print(ETS_scores)


##              ME      RMSE      MAE      MPE      MAPE
## Test set -103.0303 1547.348 1280.656 -9.574082 28.27524
```

Model 1 STL + ETS: Forecast 2011


```
#Fit and forecast STL model January 1st to February 28th 2011
ETS_model2011 <- stlf(ts_daily2010,h=59)

#Plot forecasting
autoplot(ETS_model2011) + ylab("Forecast")
```

Forecasts from STL + ETS(A,N,N)



```
#Plot model + observed data  
autoplot(ts_daily2010) +  
  autolayer(ETS_model2011, series="STL Model",PI=FALSE) +  
  ylab("Daily Demand")
```



Model 2 TBATS: Forecast 2010


The model looks like a really bad fit visually and will not be used to forecast for 2011.

```
#Fit and forecast TBATS model
TBATS_model <- tbats(ts_daily)

#forecast
TBATS_for <- forecast(TBATS_model,h=59)


#Plot forecasting
autoplot(TBATS_for) + ylab("Forecast")
```


Forecasts from BATS(0.008, {0,3}, -, -)




```
#Plot the residuals
```

```
checkresiduals(TBATS_model)
```



```
##
##  Ljung-Box test
##
## data:  Residuals from TBATS
## Q* = 105.04, df = 68, p-value = 0.002646
##
## Model df: 5.    Total lags used: 73
#Plot model + observed data
autoplot(ts_daily2010_test) +
  autolayer(TBATS_for, series="TBATS Model",PI=FALSE) +
  ylab("Daily Demand")
```



```
#Check accuracy of model
n_for <- 59
observed <- df_processed2010[1827:1885, "Daily_data"]
TBATS_scores <- accuracy(TBATS_for$mean,observed)
print(TBATS_scores)
```

```
##              ME      RMSE      MAE      MPE      MAPE
## Test set 2226.34 2499.097 2234.019 40.94677 41.24095
```

Model 2 TBATS: Forecast 2011


```
#Fit and forecast TBATS model
TBATS_model2011 <- tbats(ts_daily2010)

## Warning in tbats(ts_daily2010): Missing values encountered. Using longest
## contiguous portion of time series

#forecast
TBATS_for2011 <- forecast(TBATS_model2011,h=59)

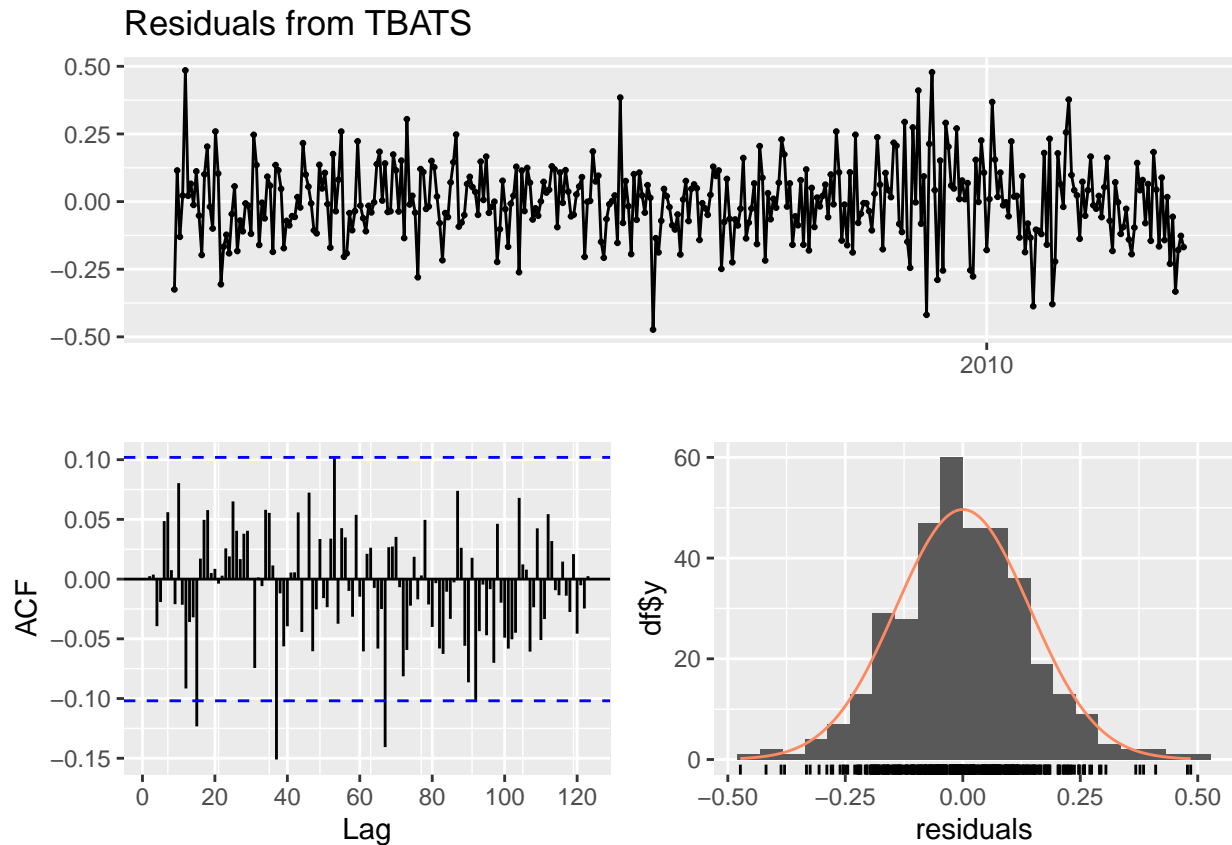
#Plot forecasting
autoplot(TBATS_for2011) + ylab("Forecast")
```

Forecasts from BATS(0.003, {2,1}, -, -)




#Plot the residuals

```
checkresiduals(TBATS_model2011)
```



```
##
##  Ljung-Box test
##
## data:  Residuals from TBATS
## Q* = 73.551, df = 69, p-value = 0.3315
##
## Model df: 5.    Total lags used: 74
#Plot model + observed data
autoplot(ts_daily2010) +
  autolayer(TBATS_for2011, series="TBATS Model",PI=FALSE) +
  ylab("Daily Demand")
```




Model 3 ARIMA + FOURIER terms: Forecast 2010

```
#Fit and forecast TBATS model
ARIMA_Four_model <- auto.arima(ts_daily,
                               seasonal=FALSE,
                               lambda=0,
                               xreg=fourier(ts_daily,
                                             K=c(2,12))
                               )


#Forecast
ARIMA_Four_for <- forecast(ARIMA_Four_model,
                           xreg=fourier(ts_daily,
                                         K=c(2,12),
                                         h=59),
                           h=59
                           )

#Plot foresting results
autoplot(ARIMA_Four_for) + ylab("Daily Demand")
```

Forecasts from Regression with ARIMA(5,1,0) errors

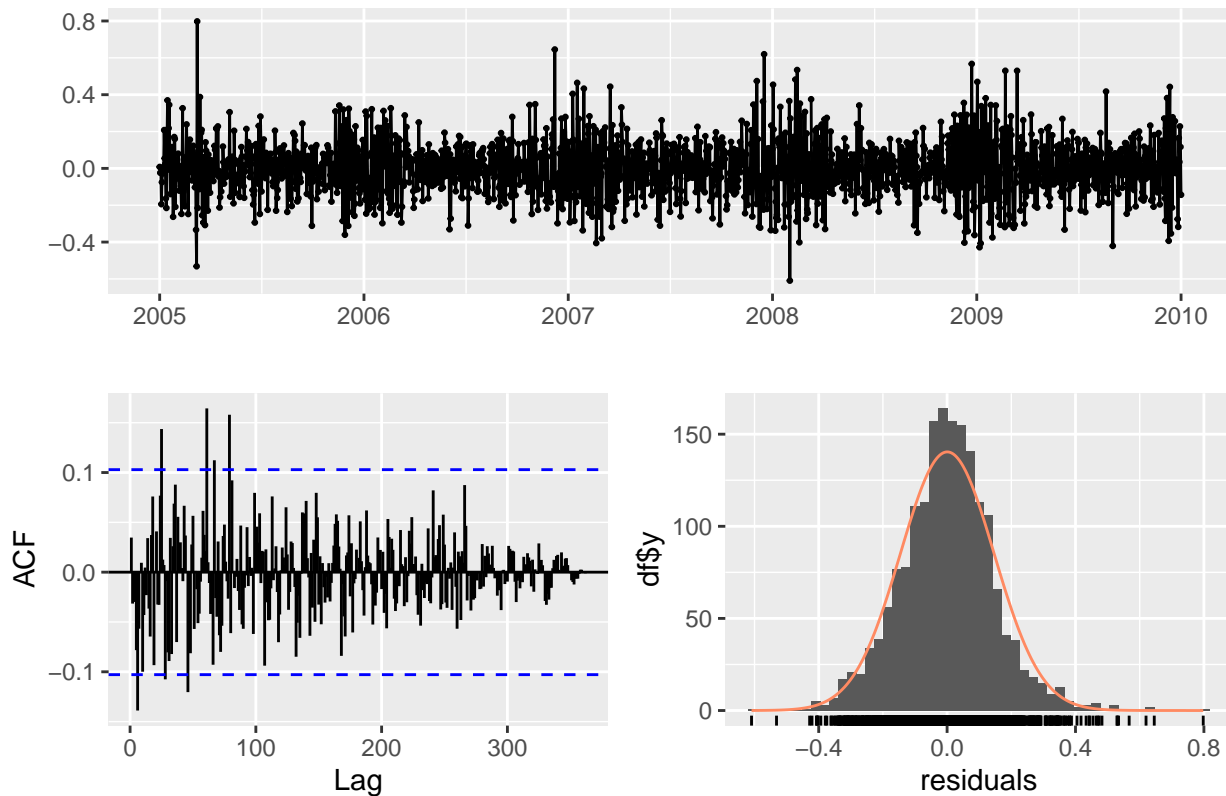


```
#Plot model + observed data  
autoplot(ts_daily2010_test) +  
  autolayer(ARIMA_Four_for, series="ARIMA_FOURIER", PI=FALSE) +  
  ylab("Daily Demand")
```



```
# Plot the residuals  
checkresiduals(ARIMA_Four_model)
```


Residuals from Regression with ARIMA(5,1,0) errors



```
##
##  Ljung-Box test
##
## data:  Residuals from Regression with ARIMA(5,1,0) errors
## Q* = 462.43, df = 360, p-value = 0.0002078
##
## Model df: 5.    Total lags used: 365
#Check accuracy of model
n_for <- 59
observed <- df_processed2010[1827:1885, "Daily_data"]
ARIMA_Four_scores <- accuracy(ARIMA_Four_for$mean,observed)
print(ARIMA_Four_scores)
```

```
##              ME      RMSE      MAE      MPE      MAPE
## Test set -284.3608 1536.398 1243.755 -13.07091 28.0451
```


Model 3 ARIMA + FOURIER terms: Forecast 2011

```
#Fit and forecast TBATS model
ARIMA_Four_model2011 <- auto.arima(ts_daily2010,
                                   seasonal=FALSE,
                                   lambda=0,
                                   xreg=fourier(ts_daily2010,
                                                K=c(2,12))
                                   )
```


```
#Forecast
ARIMA_Four_for2011 <- forecast(ARIMA_Four_model2011,
                               xreg=fourier(ts_daily2010,
                                             K=c(2,12),
                                             h=59),
                               h=59
                               )

#Plot forecasting results
autoplot(ARIMA_Four_for2011) + ylab("Daily Demand")
```

Forecasts from Regression with ARIMA(5,1,0) errors



```
#Plot model + observed data
autoplot(ts_daily2010) +
  autolayer(ARIMA_Four_for2011, series="ARIMA_FOURIER", PI=FALSE) +
  ylab("Daily Demand")
```



```
#print(ARIMA_Four_for2011$mean)
```


Model 4 Neural Network Time Series: Forecasts 2010

```
#NN_fit <- nnetar(ts_act_power_daily_train,p=1,P=1)
NN_model <- nnetar(ts_daily,decay=0.5, maxit=150, p=1,P=0,xreg=fourier(ts_daily, K=c(2,12)))


#NN_for <- forecast(NN_fit, h=365)
NN_for <- forecast(NN_model, h=59,xreg=fourier(ts_daily,
                                              K=c(2,12),h=59))

#Plot forecasting results
autoplot(NN_for) +
  ylab("Daily Demand")
```

Forecasts from NNAR(1,15)






```
#Plot model + observed data  
autoplot(ts_daily2010_test) +  
  autolayer(NN_for, series="Daily Demand",PI=FALSE)+  
  ylab("Daily Demand")
```



```
checkresiduals(NN_model)
```

Residuals from NNAR(1,15)



```
#Checking error variables to decide which model fits the data the best
n_for <- 59
observed <- df_processed2010[1827:1885, "Daily_data"]
NN_scores1 <- accuracy(NN_for$mean,observed)
print(NN_scores1)
```

```
##           ME      RMSE      MAE      MPE      MAPE
## Test set 1007.775 1924.063 1569.446 13.25849 30.17092
```

Model 4 Neural Network Time Series: Forecasts 2011


Note: Based on the error variables calculated with the accuracy() function, the Neural Network model seems to fit the data the best.

```
#NN_fit <- nnetar(ts_act_power_daily_train,p=1,P=1)
NN_model2010 <- nnetar(ts_daily2010,lambda = 0.5,p=1,P=0,xreg=fourier(ts_daily2010, K=c(2,12)))

#NN_for <- forecast(NN_fit, h=365)
NN_for2010 <- forecast(NN_model2010,PI=TRUE, h=59,xreg=fourier(ts_daily,
                                                              K=c(2,12),h=59))


#Plot forecasting results
autoplot(NN_for2010) +
  ylab("Daily Demand")
```

Forecasts from NNAR(1,15)



```
#Plot model + observed data
autoplot(ts_daily2010) +
  autolayer(NN_for2010, series="Daily_data",PI=FALSE)+
```

```
ylab("Daily Demand")
```



```
print(NN_for2010$mean)
```

```
## Multi-Seasonal Time Series:
## Start: 2011 2
## Seasonal Periods: 7 365.25
## Data:
## [1] 5110.207 5434.845 5625.082 5131.692 5098.054 5412.085 5184.238 4857.057
## [9] 4889.952 4021.525 3743.225 3829.447 4017.425 3850.678 3663.322 3855.889
## [17] 4051.004 3976.768 4057.312 4323.346 4471.939 4525.479 4761.202 4666.672
## [25] 4419.659 4098.698 3789.777 4353.527 4883.824 5002.769 4536.884 4402.893
## [33] 4663.363 4779.053 5027.869 4966.524 4681.830 4684.892 4709.388 4243.761
## [41] 3837.482 3997.712 4078.363 4118.396 4261.697 4013.714 3694.314 3533.352
## [49] 3576.141 4192.222 4422.809 4193.911 3732.190 3489.900 3875.511 4111.670
## [57] 3937.347 3784.880 4326.572
```

```
print(ARIMA_Four_for2011$mean)
```

```
## Multi-Seasonal Time Series:
## Start: 2011 2
## Seasonal Periods: 7 365.25
## Data:
## [1] 5526.415 5794.942 5726.939 5667.654 5706.920 5639.929 5716.611 5964.147
## [9] 6024.705 5911.567 5926.901 6015.781 6071.908 6284.935 6621.350 6715.763
## [17] 6611.811 6643.567 6751.663 6812.521 7027.192 7351.693 7383.301 7181.735
## [25] 7117.411 7124.783 7073.802 7174.451 7377.935 7285.460 6973.496 6809.690
## [33] 6728.359 6607.167 6643.298 6789.783 6681.093 6389.449 6250.021 6200.795
## [41] 6127.257 6210.527 6407.701 6371.016 6159.486 6090.784 6105.596 6089.920
```

Table 1: Forecast Accuracy for Seasonal Data

	ME	RMSE	MAE	MPE	MAPE
STL-ETS	-103.0303	1547.348	1280.656	-9.57408	28.27524
TBATS	2226.3400	2499.097	2234.019	40.94677	41.24095
ARIMA_FOUR	-284.3608	1536.398	1243.755	-13.07091	28.04510
NEU-NETWORK	1007.7749	1924.063	1569.446	13.25849	30.17092

```
## [49] 6222.002 6459.386 6448.920 6246.259 6173.030 6169.368 6120.337 6205.469
## [57] 6380.416 6298.339 6023.813
```

Compare performance matrix

```
#create data frame
seas_scores <- as.data.frame(rbind(ETS_scores, TBATS_scores, ARIMA_Four_scores, NN_scores1))
row.names(seas_scores) <- c("STL-ETS", "TBATS", "ARIMA_FOUR", "NEU-NETWORK")

#choose model with lowest RMSE
best_model_index <- which.min(seas_scores[, "RMSE"])
cat("The best model by RMSE is:", row.names(seas_scores[best_model_index,]))

## The best model by RMSE is: ARIMA_FOUR

kbl(seas_scores,
     caption = "Forecast Accuracy for Seasonal Data",
     digits = array(5, ncol(seas_scores))) %>%
  kable_styling(full_width = FALSE, position = "center") %>%
  #highlight model with lowest RMSE
  kable_styling(latex_options = "striped", stripe_index = which.min(seas_scores[, "RMSE"]))
```